

etl

September 23, 2023

1 ETL Processes

Use this notebook to develop the ETL process for each of your tables before completing the `etl.py` file to load the whole datasets.

```
In [1]: import os
import glob
import psycopg2
import pandas as pd
from sql_queries import *
```

```
In [2]: conn = psycopg2.connect("host=127.0.0.1 dbname=sparkifydb user=student password=student")
cur = conn.cursor()
```

```
In [3]: def get_files(filepath):
    all_files = []
    for root, dirs, files in os.walk(filepath):
        files = glob.glob(os.path.join(root, '*.json'))
        for f in files :
            all_files.append(os.path.abspath(f))

    return all_files
```

2 Process song_data

In this first part, you'll perform ETL on the first dataset, `song_data`, to create the songs and artists dimensional tables.

Let's perform ETL on a single song file and load a single record into each table to start. - Use the `get_files` function provided above to get a list of all song JSON files in `data/song_data` - Select the first song in this list - Read the song file and view the data

```
In [4]: song_files = get_files("data/song_data")
```

```
In [5]: filepath = song_files[0]
```

```
In [6]: df = pd.read_json(filepath, lines=True)
df.head()
```

```

Out[6]:
      artist_id  artist_latitude  artist_location  artist_longitude \
0  ARD7TVE1187B99BFB1          NaN  California - LA          NaN

      artist_name  duration  num_songs  song_id  title \
0      Casual    218.93179          1  SOMZWCG12A8C13C480  I Didn't Mean To

      year
0      0

```

2.1 #1: songs Table

Extract Data for Songs Table

- Select columns for song ID, title, artist ID, year, and duration
- Use `df.values` to select just the values from the dataframe
- Index to select the first (only) record in the dataframe
- Convert the array to a list and set it to `song_data`

```

In [7]: selected_columns = ['song_id', 'title', 'artist_id', 'year', 'duration']
        song_data = df[selected_columns].values[0].tolist()
        song_data

```

```

Out[7]: ['SOMZWCG12A8C13C480', 'I Didn't Mean To', 'ARD7TVE1187B99BFB1', 0, 218.93179]

```

Insert Record into Song Table Implement the `song_table_insert` query in `sql_queries.py` and run the cell below to insert a record for this song into the songs table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the songs table in the sparkify database.

```

In [8]: cur.execute(song_table_insert, song_data)
        conn.commit()

```

Run `test.ipynb` to see if you've successfully added a record to this table.

2.2 #2: artists Table

Extract Data for Artists Table

- Select columns for artist ID, name, location, latitude, and longitude
- Use `df.values` to select just the values from the dataframe
- Index to select the first (only) record in the dataframe
- Convert the array to a list and set it to `artist_data`

```

In [9]: selected_columns1 = ['artist_id', 'artist_name', 'artist_location', 'artist_latitude', '
        artist_data = df[selected_columns1].values[0].tolist()
        artist_data

```

```

Out[9]: ['ARD7TVE1187B99BFB1', 'Casual', 'California - LA', nan, nan]

```

Insert Record into Artist Table Implement the `artist_table_insert` query in `sql_queries.py` and run the cell below to insert a record for this song's artist into the artists table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the artists table in the sparkify database.

```
In [10]: cur.execute(artist_table_insert, artist_data)
         conn.commit()
```

Run `test.ipynb` to see if you've successfully added a record to this table.

3 Process log_data

In this part, you'll perform ETL on the second dataset, `log_data`, to create the time and users dimensional tables, as well as the songplays fact table.

Let's perform ETL on a single log file and load a single record into each table. - Use the `get_files` function provided above to get a list of all log JSON files in `data/log_data` - Select the first log file in this list - Read the log file and view the data

```
In [11]: log_files = get_files("data/log_data")
```

```
In [12]: filepath = log_files[0]
```

```
In [13]: df = pd.read_json(filepath, lines=True)
         df.head()
```

```
Out[13]:
```

	artist	auth	firstName	\
0	None	Logged In	Celeste	
1	Pavement	Logged In	Sylvie	
2	Barry Tuckwell/Academy of St Martin-in-the-Fie...	Logged In	Celeste	
3	Gary Allan	Logged In	Celeste	
4	None	Logged In	Jacqueline	

	gender	itemInSession	lastName	length	level	\
0	F	0	Williams	NaN	free	
1	F	0	Cruz	99.16036	free	
2	F	1	Williams	277.15873	free	
3	F	2	Williams	211.22567	free	
4	F	0	Lynch	NaN	paid	

	location	method	page	\
0	Klamath Falls, OR	GET	Home	
1	Washington-Arlington-Alexandria, DC-VA-MD-WV	PUT	NextSong	
2	Klamath Falls, OR	PUT	NextSong	
3	Klamath Falls, OR	PUT	NextSong	
4	Atlanta-Sandy Springs-Roswell, GA	GET	Home	

	registration	sessionId	song	\
0	1.541078e+12	438	None	

1	1.540266e+12	345	Mercy:The Laundromat
2	1.541078e+12	438	Horn Concerto No. 4 in E flat K495: II. Romanc...
3	1.541078e+12	438	Nothing On But The Radio
4	1.540224e+12	389	None

	status	ts	userAgent \
0	200	1541990217796	"Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit...
1	200	1541990258796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...
2	200	1541990264796	"Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit...
3	200	1541990541796	"Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit...
4	200	1541990714796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...

	userId
0	53
1	10
2	53
3	53
4	29

3.1 #3: time Table

Extract Data for Time Table

- Filter records by NextSong action
- Convert the ts timestamp column to datetime
- Hint: the current timestamp is in milliseconds
- Extract the timestamp, hour, day, week of year, month, year, and weekday from the ts column and set time_data to a list containing these values in order
- Hint: use pandas' [dt attribute](#) to access easily datetimelike properties.
- Specify labels for these columns and set to column_labels
- Create a dataframe, time_df, containing the time data for this file by combining column_labels and time_data into a dictionary and converting this into a dataframe

```
In [14]: df = df.loc[df['page'] == 'NextSong']
df.head()
```

```
Out[14]:
```

	artist	auth	firstName \
1	Pavement	Logged In	Sylvie
2	Barry Tuckwell/Academy of St Martin-in-the-Fie...	Logged In	Celeste
3	Gary Allan	Logged In	Celeste
5	Charttraxx Karaoke	Logged In	Celeste
6	The Libertines	Logged In	Jacqueline

	gender	itemInSession	lastName	length	level \
1	F	0	Cruz	99.16036	free
2	F	1	Williams	277.15873	free
3	F	2	Williams	211.22567	free
5	F	3	Williams	225.17506	free
6	F	1	Lynch	179.53914	paid

	location	method	page	\
1	Washington-Arlington-Alexandria, DC-VA-MD-WV	PUT	NextSong	
2	Klamath Falls, OR	PUT	NextSong	
3	Klamath Falls, OR	PUT	NextSong	
5	Klamath Falls, OR	PUT	NextSong	
6	Atlanta-Sandy Springs-Roswell, GA	PUT	NextSong	

	registration	sessionId	song	\
1	1.540266e+12	345	Mercy:The Laundromat	
2	1.541078e+12	438	Horn Concerto No. 4 in E flat K495: II. Romanc...	
3	1.541078e+12	438	Nothing On But The Radio	
5	1.541078e+12	438	Fireflies	
6	1.540224e+12	389	The Good Old Days	

	status	ts	userAgent	\
1	200	1541990258796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	
2	200	1541990264796	"Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit...	
3	200	1541990541796	"Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit...	
5	200	1541990752796	"Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit...	
6	200	1541990842796	"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4...	

	userId
1	10
2	53
3	53
5	53
6	29

```
In [15]: t = pd.to_datetime(df['ts'], unit='ms')
t.head()
```

```
Out[15]: 1    2018-11-12 02:37:38.796
2    2018-11-12 02:37:44.796
3    2018-11-12 02:42:21.796
5    2018-11-12 02:45:52.796
6    2018-11-12 02:47:22.796
Name: ts, dtype: datetime64[ns]
```

```
In [16]: time_data = ([
            df.ts.values,
            t.dt.hour.values,
            t.dt.day.values,
            t.dt.weekofyear.values,
            t.dt.month.values,
            t.dt.year.values,
            t.dt.weekday.values,
        ])
column_labels = ('time_stamp', 'hour', 'day', 'week', 'month', 'year', 'weekday')
```

```
In [17]: time_df = pd.DataFrame(dict(zip(column_labels, time_data)))
        time_df.head()
```

```
Out[17]:
```

	time_stamp	hour	day	week	month	year	weekday
0	1541990258796	2	12	46	11	2018	0
1	1541990264796	2	12	46	11	2018	0
2	1541990541796	2	12	46	11	2018	0
3	1541990752796	2	12	46	11	2018	0
4	1541990842796	2	12	46	11	2018	0

Insert Records into Time Table Implement the `time_table_insert` query in `sql_queries.py` and run the cell below to insert records for the timestamps in this log file into the time table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the time table in the sparkify database.

```
In [18]: for i, row in time_df.iterrows():
        cur.execute(time_table_insert, list(row))
        conn.commit()
```

Run `test.ipynb` to see if you've successfully added records to this table.

3.2 #4: users Table

Extract Data for Users Table

- Select columns for user ID, first name, last name, gender and level and set to `user_df`

```
In [19]: user_df = df[['userId', 'firstName', 'lastName', 'gender', 'level']]
```

Insert Records into Users Table Implement the `user_table_insert` query in `sql_queries.py` and run the cell below to insert records for the users in this log file into the users table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the users table in the sparkify database.

```
In [20]: for i, row in user_df.iterrows():
        cur.execute(user_table_insert, row)
        conn.commit()
```

Run `test.ipynb` to see if you've successfully added records to this table.

3.3 #5: songplays Table

Extract Data and Songplays Table This one is a little more complicated since information from the songs table, artists table, and original log file are all needed for the songplays table. Since the log file does not specify an ID for either the song or the artist, you'll need to get the song ID and artist ID by querying the songs and artists tables to find matches based on song title, artist name, and song duration time. - Implement the `song_select` query in `sql_queries.py` to find the song ID and artist ID based on the title, artist name, and duration of a song. - Select the timestamp, user ID, level, song ID, artist ID, session ID, location, and user agent and set to `songplay_data`

Insert Records into Songplays Table

- Implement the `songplay_table_insert` query and run the cell below to insert records for the songplay actions in this log file into the `songplays` table. Remember to run `create_tables.py` before running the cell below to ensure you've created/resetted the `songplays` table in the sparkify database.

```
In [21]: for index, row in df.iterrows():

        # get songid and artistid from song and artist tables
        cur.execute(song_select, (row.song, row.artist, row.length))
        results = cur.fetchone()

        if results:
            songid, artistid = results
        else:
            songid, artistid = None, None

        # insert songplay record
        songplay_data = (row.ts, row.userId, row.level, songid, artistid, row.sessionId, row.location)
        cur.execute(songplay_table_insert, songplay_data)
        conn.commit()
```

Run `test.ipynb` to see if you've successfully added records to this table.

4 Close Connection to Sparkify Database

```
In [22]: conn.close()
```

5 Implement `etl.py`

Use what you've completed in this notebook to implement `etl.py`.

```
In [ ]:
```